

Advanced KEM Concepts

(Hybrid) Obfuscation and Verifiable Decapsulation

IBM Research Security

Felix Günther

IBM Research Europe – Zurich

based on work with

Lewis Glabush

EPFL

Douglas Stebila

U Waterloo

Britta Hale

NPS

Shannon Veitch

ETH Zurich

Kathrin Hövelmanns

TU Eindhoven

Elizabeth van Oorschot

McGill

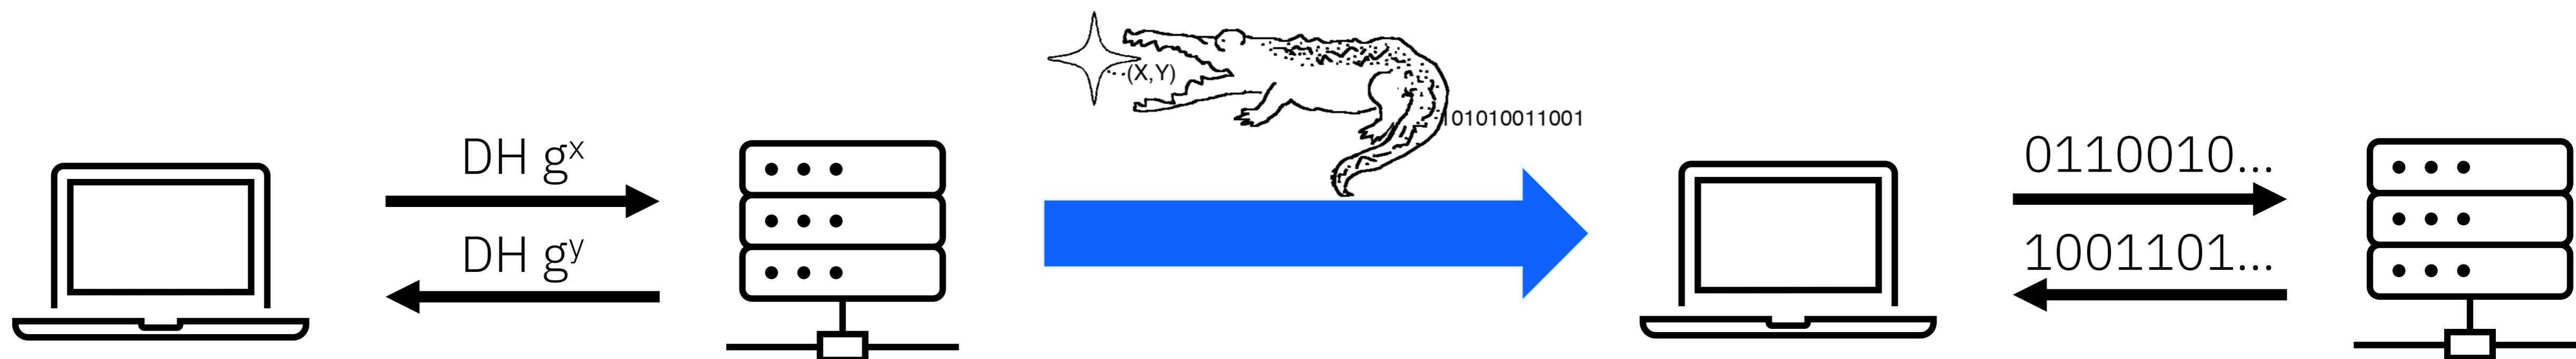
Michael Rosenberg

Cloudflare

Protocol Obfuscation

Internet protocols hide **metadata** to protect user privacy, dissuade protocol fingerprinting, and prevent network ossification

- TLS 1.3 Encrypted Client Hello, QUIC, obfs4, Shadowsocks, ...
- “Fully-encrypted” protocols, with **obfuscated** key exchange

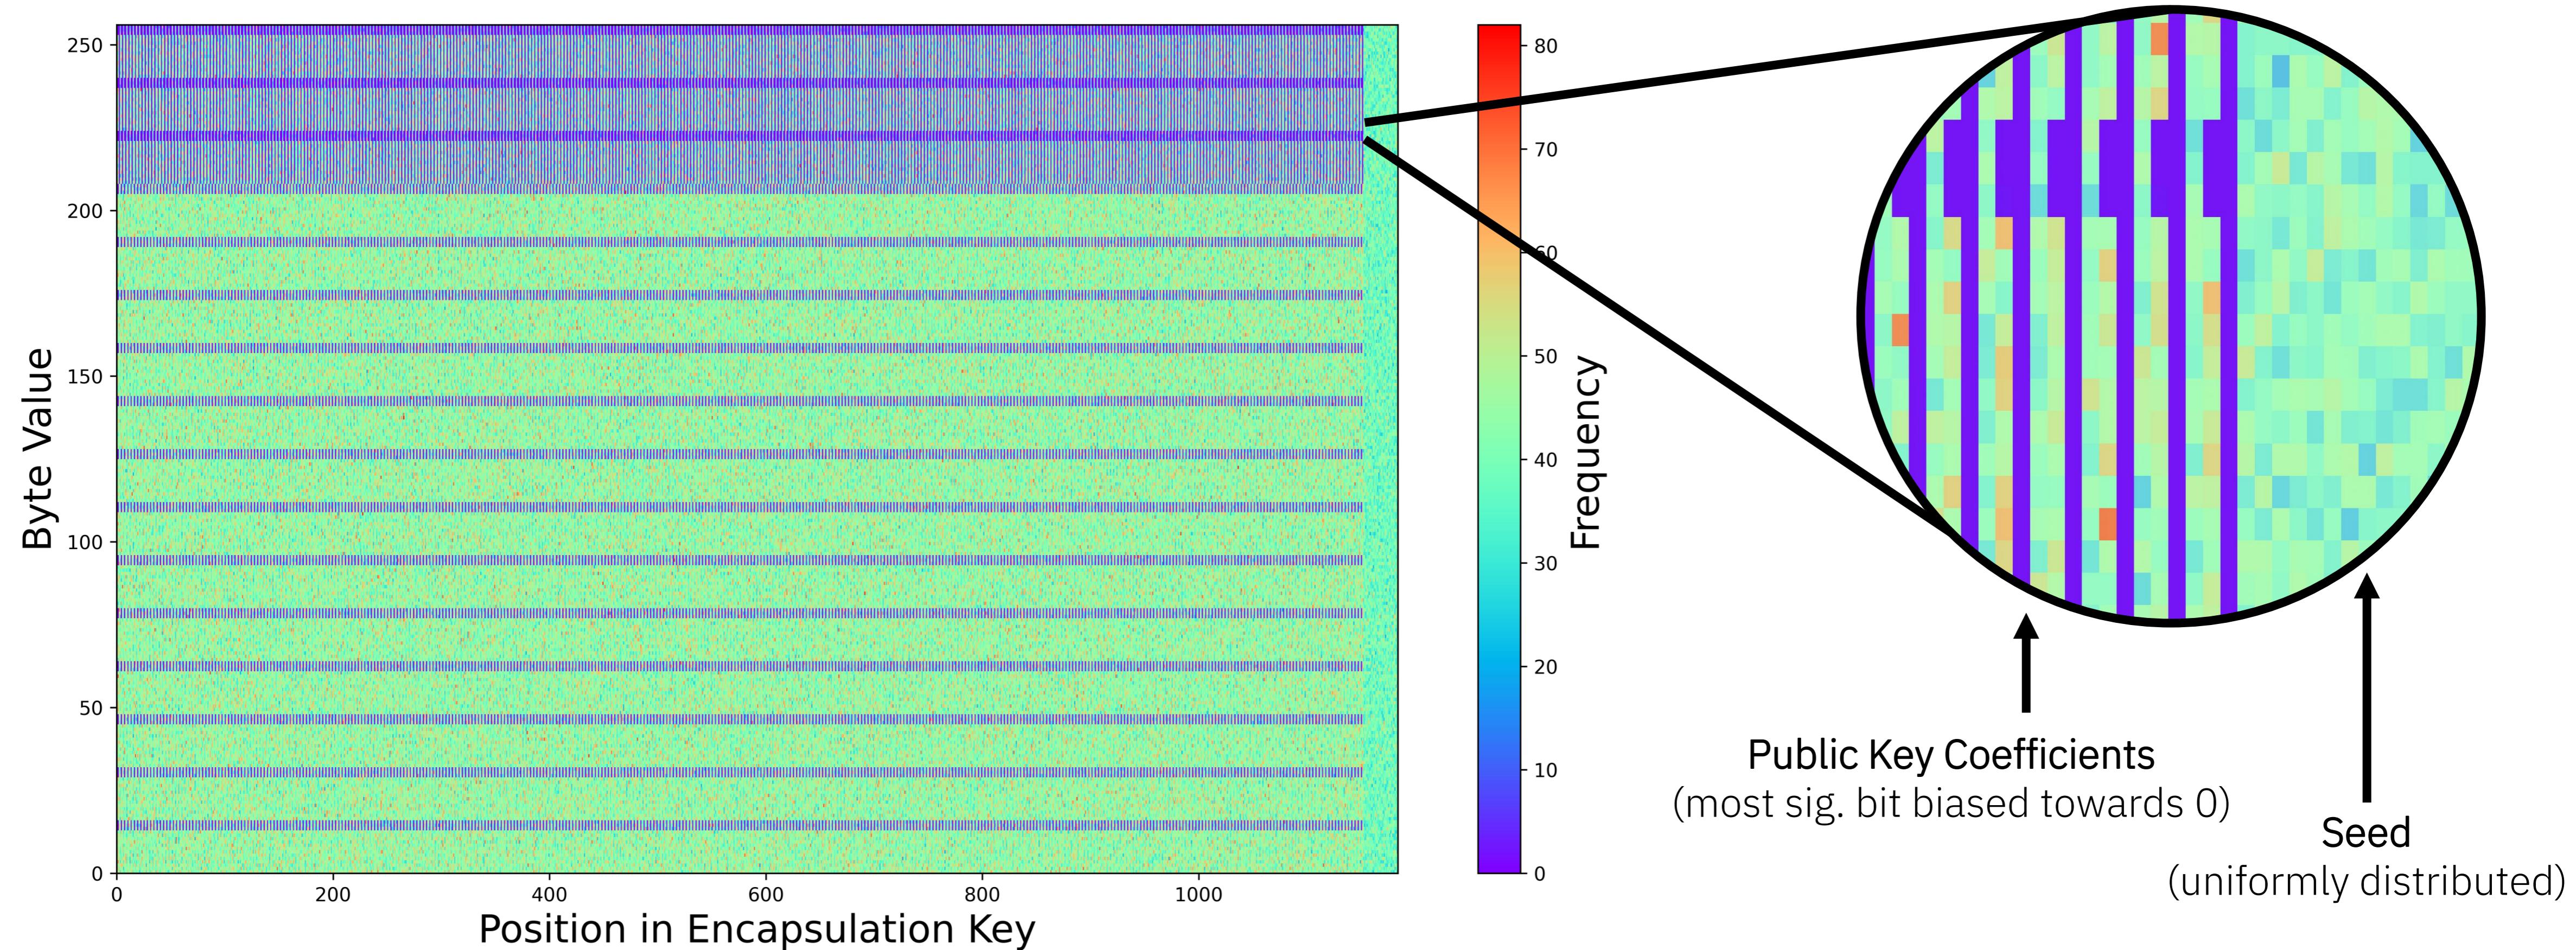


Quantum-safe transition?

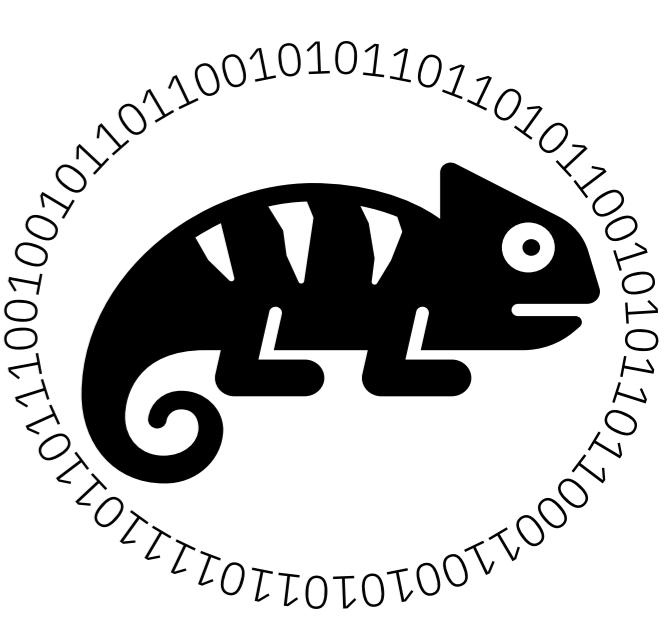
ML-KEM public keys and ciphertexts don't look random!

Byte Distribution of ML-KEM-768 Public Keys

ML-KEM public keys: $t = As + e$ (vectors of coefficients mod $q = 3329$), ρ (random seed generating A)



Kemeleon



ML-KEM public keys

- vector of coefficients mod q = 3329

$[a_1] [a_2] [a_3] \dots [a_b] \quad a_i \in \mathbb{Z}_q = \{0, \dots, 3328\}$ - each a_i represented in 12 bits

most sig. bit of each value biased towards 0

- Encoding for public keys:
 - 1. accumulate into one big number
 - 2. rejection sampling: reject if msb is 1

$$[A = a_1 + a_2 \cdot q + a_3 \cdot q^2 + \cdots + a_b \cdot q^{b-1}]$$

↑

most sig. bit still biased towards 0

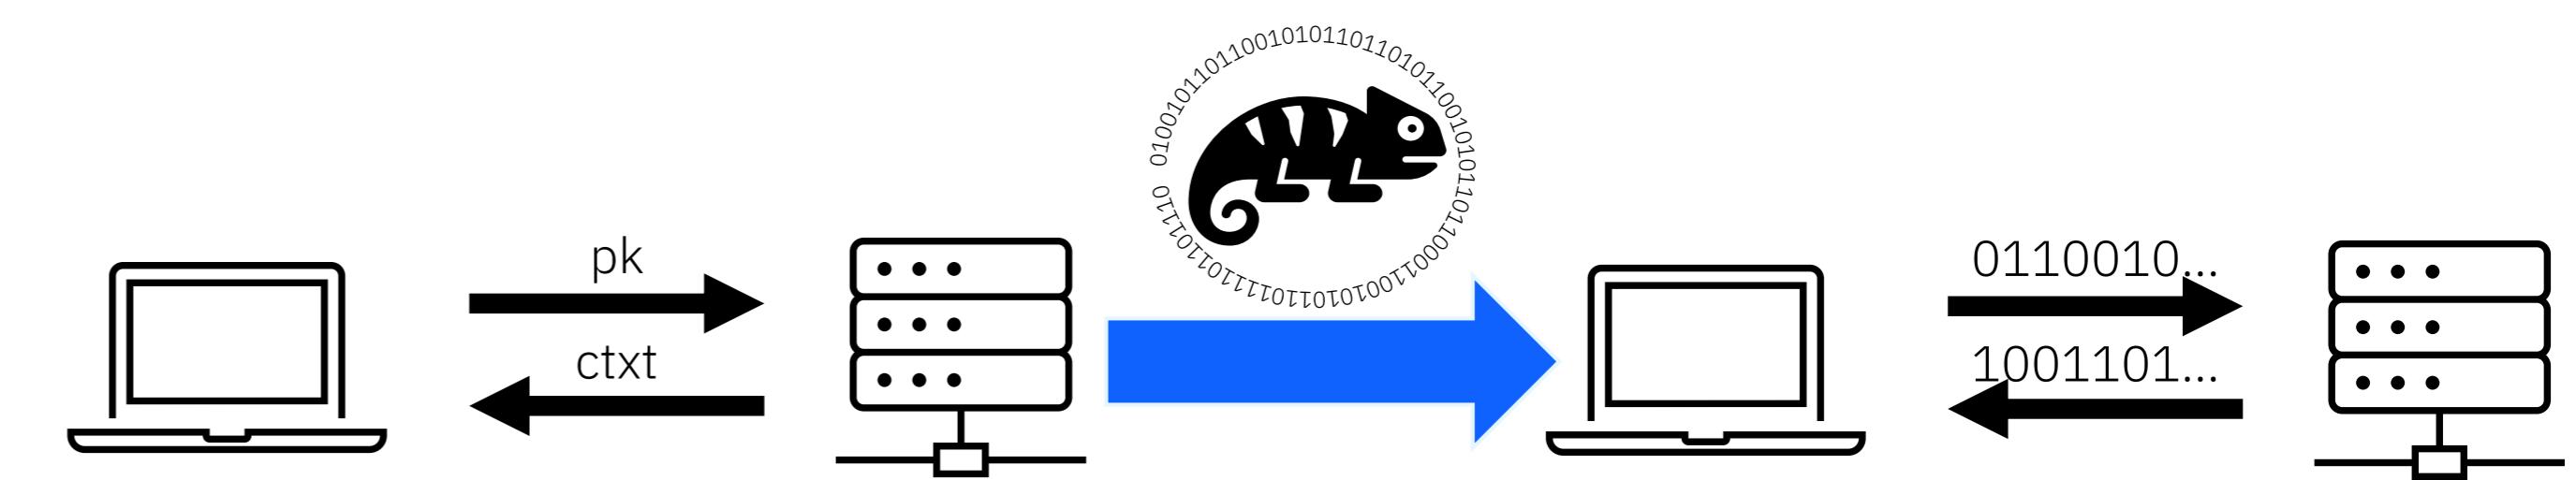
Encoded public keys ~2.5% smaller than regular
(-19/28/38 bytes for ML-KEM-512/768/1024)

ML-KEM-768 likelihood of rejection is ~17%

ML-KEM ciphertexts

- vector of **compressed** coefficients – need to first “decompress”
 - encoded ciphertexts larger than regular (6–15%)

Obfuscated KEMs

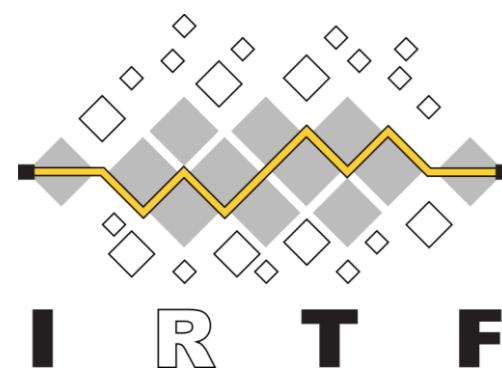


ML-KEM

+ Kemeleon public key and ciphertext encoding

= Obfuscated KEM: **ML-Kemeleon**

- IND-CCA: indistinguishability of shared secrets
- SPR-CCA: ind. of secrets + ciphertexts simulatable (implies **anonymity**)
- Ciphertext/Public-key Uniformity: indistinguishable from random bit strings



Kemeleon adopted by CFRG

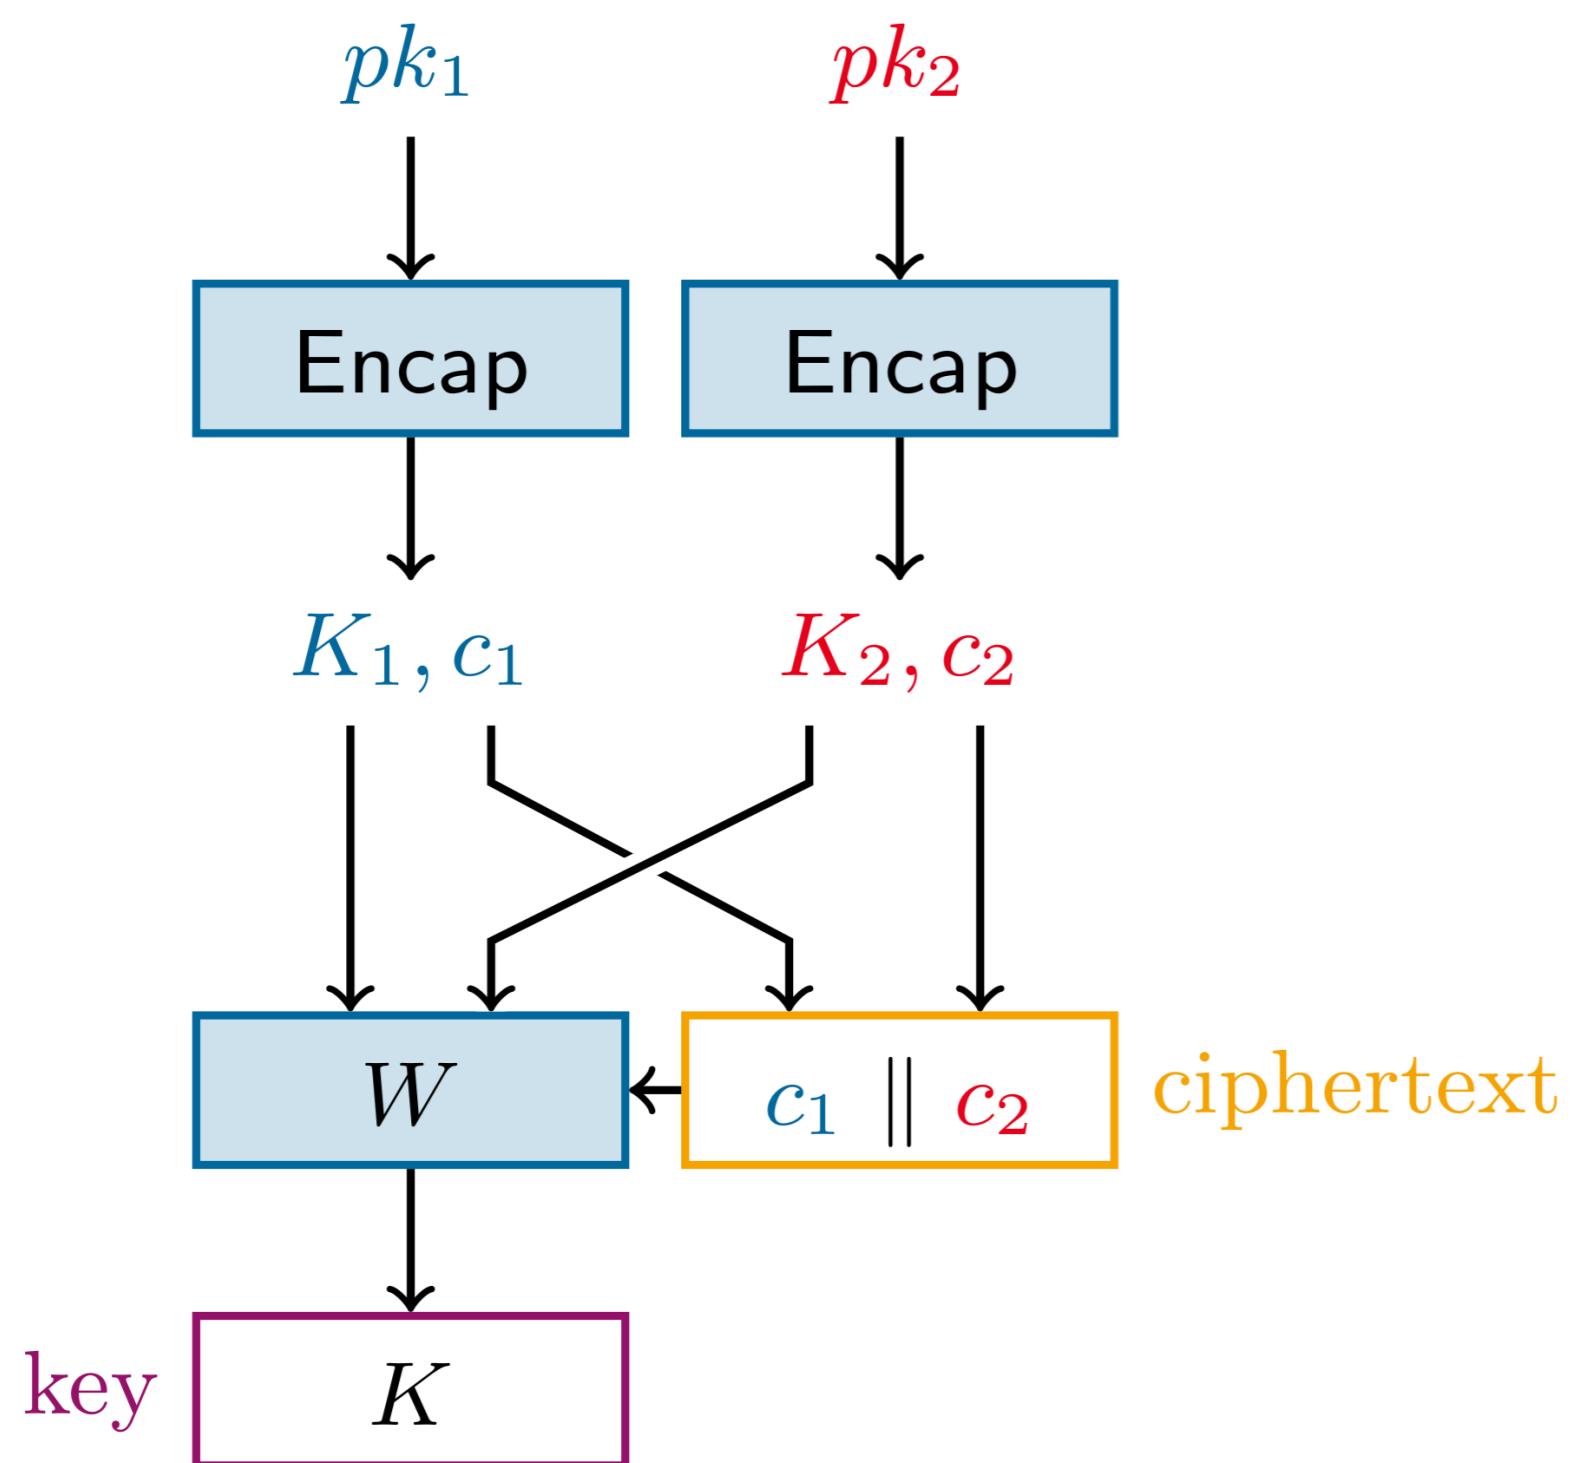
<https://datatracker.ietf.org/doc/draft-irtf-cfrg-kemeleon/>

(more variants: no rejection, deterministic, ...)

Hybrid KEMs

Parallel Combiner

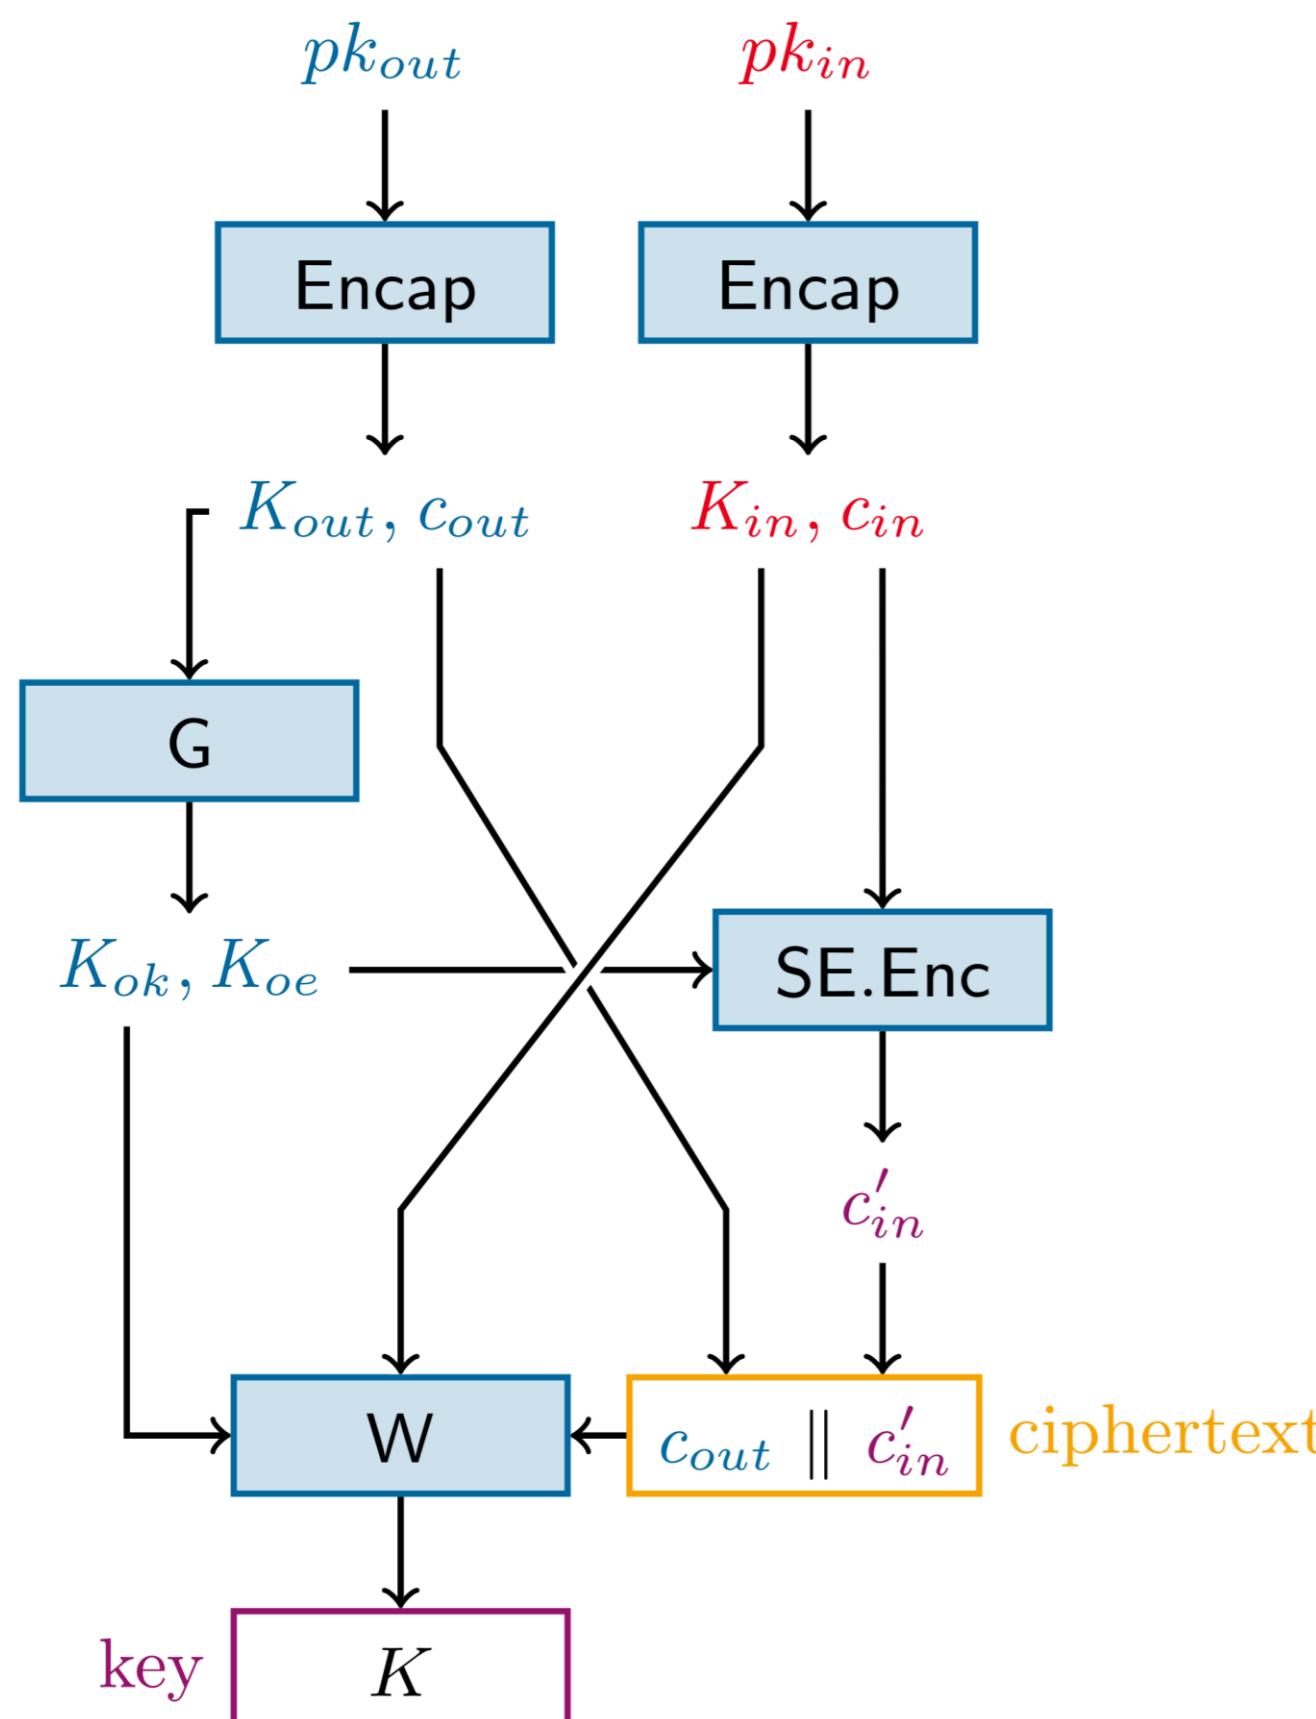
TLS 1.3 hybrid, HPKE Xyber, XWing, QSF, KitchenSink, Chempat, ...



- Hybrid IND-CCA security
- Hybrid Obfuscation

Hybrid Obfuscated KEMs

OEINC



Outer-encrypts-inner nested combiner

- Hybrid IND-CCA security
- Hybrid Obfuscation
- Low overhead: 1 PRG + 1 XOR

example: outer = DH-Elligator
inner = ML-Kemeleon

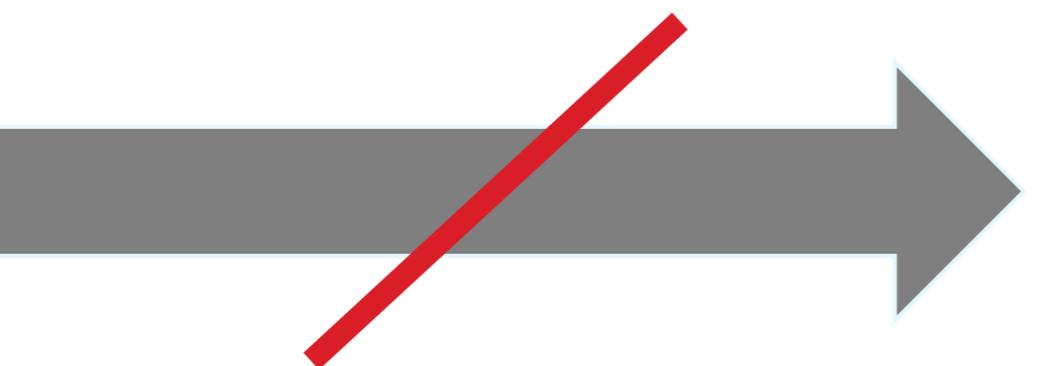
(statistical)
(computational)

Use OEINC to build

- hybrid obfuscated key exchange
- hybrid PAKE (w/ adaptive corruptions)

Cryptography Is Brittle

functionality

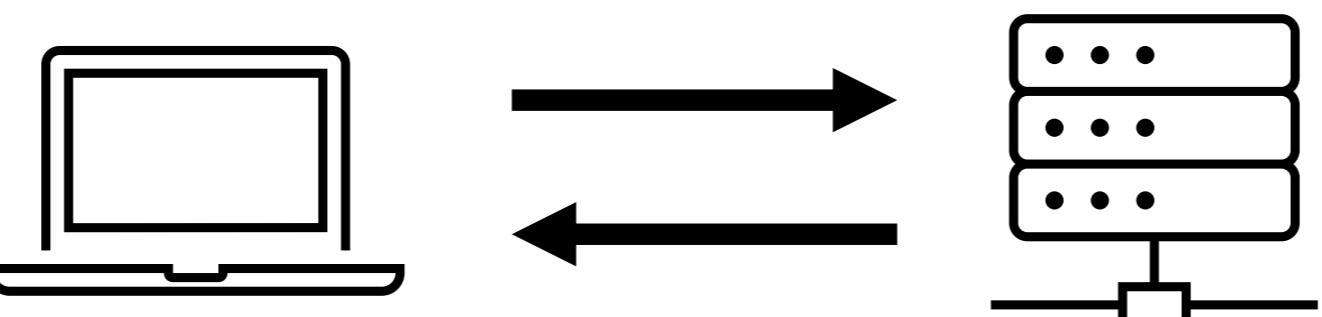


security

```
static OSStatus
SSLVerifySignedServerKeyExchange(SSLContext *ct
    SSLBuffer signedParams, uint8_t *signature,
{
    OSStatus      err;
    ...

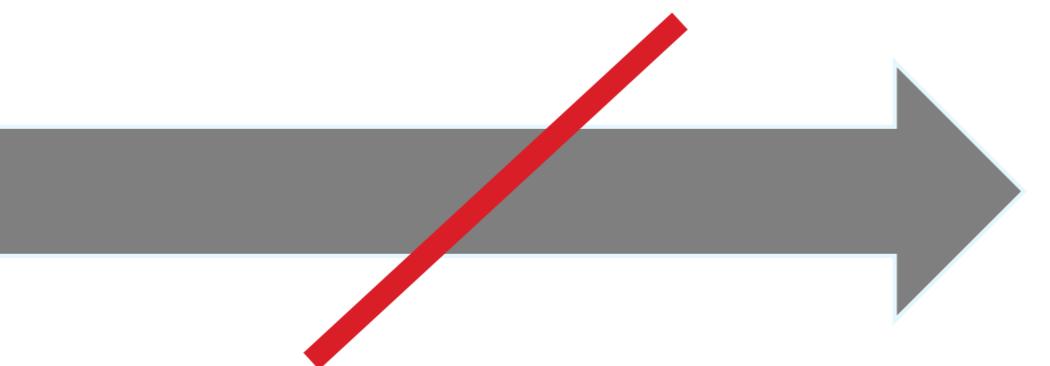
    if ((err = SSLHashSHA1.update(&hashCtx, &se
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &si
        goto fail;
        goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &has
        goto fail;
    ...

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}
```



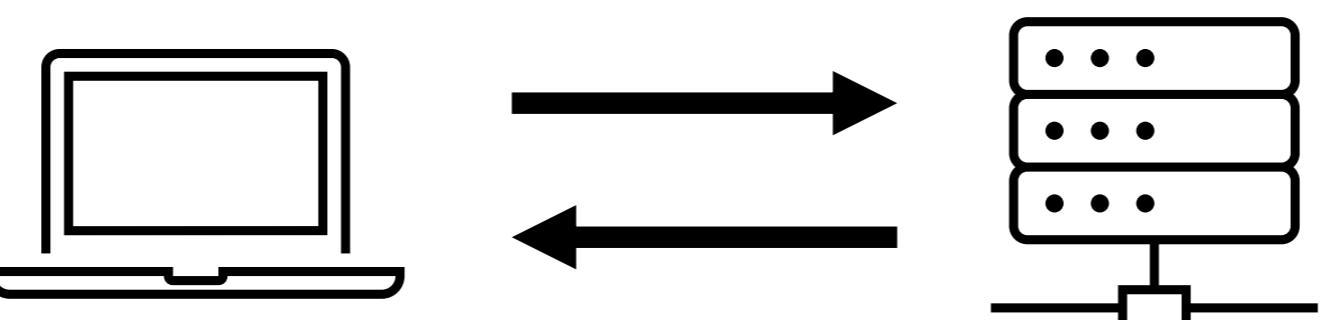
Cryptography Is Brittle

functionality



security

```
static OSStatus  
SSLVerifySignedServerKeyExchange(SSLContext *ct  
    SSLBuffer signedParams, uint8_t *signature,  
{  
    OSStatus err;  
    ...  
    if ((err = SSLHashSHA1.update(&hashCtx, &se  
        goto fail;  
    if ((err = SSLHashSHA1.update(&hashCtx, &si  
        goto fail;  
        goto fail;  
    if ((err = SSLHashSHA1.final(&hashCtx, &has  
        goto fail;  
    ...  
  
fail:  
    SSLFreeBuffer(&signedHashes);  
    SSLFreeBuffer(&hashCtx);  
    return err;  
}
```



Algorithm 18 ML-KEM.Decaps_internal(dk, c)

```
5:  $m' \leftarrow K\text{-PKE.Decrypt}(dk_{\text{PKE}}, c)$   
6:  $(K', r') \leftarrow G(m' \| h)$   
7:  $\bar{K} \leftarrow J(z \| c)$   
8:  $c' \leftarrow K\text{-PKE.Encrypt}(ek_{\text{PKE}}, m', r')$  ??  
9: if  $c \neq c'$  then  
10:    $K' \leftarrow \bar{K}$   
11: end if  
12: return  $K'$ 
```



FO transform

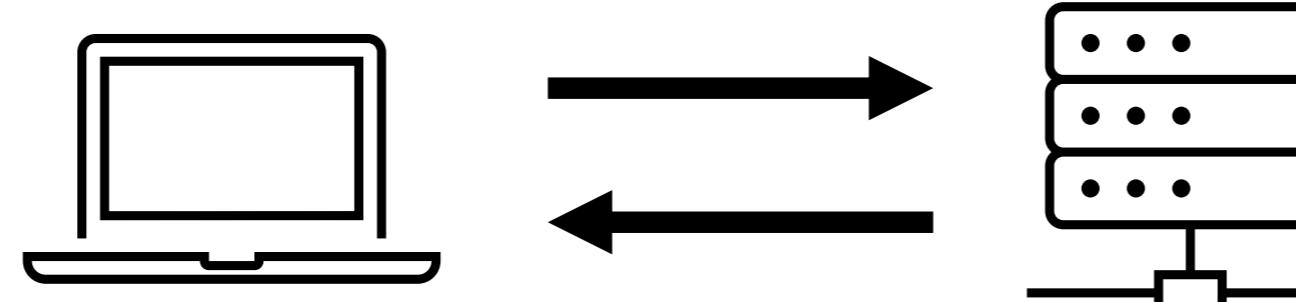
Cryptography Is Brittle

Can we tie security to basic functionality?

functionality

security

```
static OSStatus  
SSLVerifySignedServerKeyExchange(SSLContext *ct  
    SSLBuffer signedParams, uint8_t *signature,  
{  
    OSStatus err;  
    ...  
  
    if ((err = SSLHashSHA1.update(&hashCtx, &se  
        goto fail;  
    if ((err = SSLHashSHA1.update(&hashCtx, &si  
        goto fail;  
        goto fail;  
    if ((err = SSLHashSHA1.final(&hashCtx, &has  
        goto fail;  
    ...  
  
fail:  
    SSLFreeBuffer(&signedHashes);  
    SSLFreeBuffer(&hashCtx);  
    return err;  
}
```



```
85 int PQCLEAN_HQC128_CLEAN_crypto_kem_dec(uint8_t *ss, const uint8_t *ct, const uint8_t *sk) {  
86  
87     uint8_t result;  
88     uint64_t u[VEC_N_SIZE_64] = {0};  
89     uint64_t v[VEC_N1N2_SIZE_64] = {0};  
90     const uint8_t *pk = sk + SEED_BYTES;  
91     uint8_t sigma[VEC_K_SIZE_BYTES] = {0};  
92     uint8_t theta[SHAKE256_512_BYTES] = {0};  
93     uint64_t u2[VEC_N_SIZE_64] = {0};  
94     uint64_t v2[VEC_N1N2_SIZE_64] = {0};  
95     uint8_t mc[VEC_K_SIZE_BYTES + VEC_N_SIZE_BYTES + VEC_N1N2_SIZE_BYTES] = {0};  
96     uint8_t tmp[VEC_K_SIZE_BYTES + PUBLIC_KEY_BYTES + SALT_SIZE_BYTES] = {0};  
97     uint8_t *m = tmp;  
98     uint8_t *salt = tmp + VEC_K_SIZE_BYTES + PUBLIC_KEY_BYTES;  
99     shake256incctx shake256state;  
100  
101    // Retrieving u, v and d from ciphertext  
102    PQCLEAN_HQC128_CLEAN_hqc_ciphertext_from_string(u, v, salt, ct);  
103  
104    // Decrypting  
105    result = PQCLEAN_HQC128_CLEAN_hqc_pke_decrypt(m, sigma, u, v, sk);  
106  
107    // Computing theta  
108    memcpy(tmp + VEC_K_SIZE_BYTES, pk, PUBLIC_KEY_BYTES);  
109    PQCLEAN_HQC128_CLEAN_shake256_512_ds(&shake256state, theta, tmp, VEC_K_SIZE_BYTES + PUBLIC_  
110  
111    // Encrypting m'  
112    PQCLEAN_HQC128_CLEAN_hqc_pke_encrypt(u2, v2, m, theta, pk);  
113  
114    // Check if c != c'  
115    result |= PQCLEAN_HQC128_CLEAN_vect_compare((uint8_t *)u, (uint8_t *)u2, VEC_N_SIZE_BYTES);  
116    result |= PQCLEAN_HQC128_CLEAN_vect_compare((uint8_t *)v, (uint8_t *)v2, VEC_N1N2_SIZE_BYT  
117  
118    result = (uint8_t) (-((int16_t) result) >> 15);  
119  
120    for (size_t i = 0; i < VEC_K_SIZE_BYTES; ++i) {  
121        mc[i] = (m[i] & result) ^ (sigma[i] & ~result);  
122    }
```

Verifiable Decapsulation

Decaps(sk, c)

05 $m' \leftarrow \text{Dec}(\text{sk}, c)$

06 $(c' \quad) \leftarrow \text{Enc}(\text{pk}, m')$

07 check $c' = c$

08 $K' \leftarrow \text{KDF}(m', \text{pk} \quad)$

09 **return** K'

Verifiable Decapsulation

Enter: Confirmation Codes

building on ideas from Heninger, [Fischlin-G'23]

```
Decaps( $\text{sk}$ ,  $c$ )
  05  $m' \leftarrow \text{Dec}(\text{sk}, c)$ 
  06  $(c', \text{cd}') \leftarrow \text{Enc}(\text{pk}, m')$ 
  07 check  $c' = c$ 
  08  $K' \leftarrow \text{KDF}(m', \text{pk}, \text{cd}')$ 
  09 return  $K'$ 
```

Idea: **faulty implementation** of re-encryption → noticeable KEM **correctness failure**

ML-KEM with Confirmation Codes

ML-KEM ciphertext compression → lost entropy



leverage lost entropy for confirmation code

Using **12-20 bytes** of confirmation code

detect **faulty re-encryption** in ML-KM-512/768/1024

by **single test** w/ probability **~1/3**

at **≤ 3.4%** performance overhead

Algorithm 14 $\text{K-PKE}.\text{Encrypt}(\text{ek}_{\text{PKE}}, m, r)$

Uses the encryption key to encrypt a plaintext message using the randomness r .

Input: encryption key $\text{ek}_{\text{PKE}} \in \mathbb{B}^{384k+32}$.

Input: message $m \in \mathbb{B}^{32}$.

Input: randomness $r \in \mathbb{B}^{32}$.

Output: ciphertext $c \in \mathbb{B}^{32(d_u k + d_v)}$.

```
1:  $N \leftarrow 0$ 
2:  $\hat{\mathbf{t}} \leftarrow \text{ByteDecode}_{12}(\text{ek}_{\text{PKE}}[0 : 384k])$   $\triangleright$  run  $\text{ByteDecode}_{12}$   $k$  times to decode  $\hat{\mathbf{t}} \in (\mathbb{Z}_q^{256})^k$ 
3:  $\rho \leftarrow \text{ek}_{\text{PKE}}[384k : 384k + 32]$   $\triangleright$  extract 32-byte seed from  $\text{ek}_{\text{PKE}}$ 
4: for ( $i \leftarrow 0; i < k; i++$ )  $\triangleright$  re-generate matrix  $\hat{\mathbf{A}} \in (\mathbb{Z}_q^{256})^{k \times k}$  sampled in Alg. 13
5:   for ( $j \leftarrow 0; j < k; j++$ )
6:      $\hat{\mathbf{A}}[i, j] \leftarrow \text{SampleNTT}(\rho \| j \| i)$   $\triangleright j$  and  $i$  are bytes 33 and 34 of the input
7:   end for
8: end for
9: for ( $i \leftarrow 0; i < k; i++$ )  $\triangleright$  generate  $\mathbf{y} \in (\mathbb{Z}_q^{256})^k$ 
10:   $\mathbf{y}[i] \leftarrow \text{SamplePolyCBD}_{\eta_1}(\text{PRF}_{\eta_1}(r, N))$   $\triangleright \mathbf{y}[i] \in \mathbb{Z}_q^{256}$  sampled from CBD
11:   $N \leftarrow N + 1$ 
12: end for
13: for ( $i \leftarrow 0; i < k; i++$ )  $\triangleright$  generate  $\mathbf{e}_1 \in (\mathbb{Z}_q^{256})^k$ 
14:   $\mathbf{e}_1[i] \leftarrow \text{SamplePolyCBD}_{\eta_2}(\text{PRF}_{\eta_2}(r, N))$   $\triangleright \mathbf{e}_1[i] \in \mathbb{Z}_q^{256}$  sampled from CBD
15:   $N \leftarrow N + 1$ 
16: end for
17:  $e_2 \leftarrow \text{SamplePolyCBD}_{\eta_2}(\text{PRF}_{\eta_2}(r, N))$   $\triangleright$  sample  $e_2 \in \mathbb{Z}_q^{256}$  from CBD
18:  $\hat{\mathbf{y}} \leftarrow \text{NTT}(\mathbf{y})$   $\triangleright$  run  $\text{NTT}$   $k$  times
19:  $\mathbf{u} \leftarrow \text{NTT}^{-1}(\hat{\mathbf{A}}^\top \circ \hat{\mathbf{y}}) + \mathbf{e}_1$   $\triangleright$  run  $\text{NTT}^{-1}$   $k$  times
20:  $\mu \leftarrow \text{Decompress}_{\eta_1}(\text{ByteDecode}_1(m))$ 
21:  $v \leftarrow \text{NTT}^{-1}(\hat{\mathbf{t}}^\top \circ \hat{\mathbf{y}}) + e_2 + \mu$   $\triangleright$  encode plaintext  $m$  into polynomial  $v$ 
22:  $c_1 \leftarrow \text{ByteEncode}_{d_u}(\text{Compress}_{d_u}(\mathbf{u}))$   $\triangleright$  run  $\text{ByteEncode}_{d_u}$  and  $\text{Compress}_{d_u}$   $k$  times
23:  $c_2 \leftarrow \text{ByteEncode}_{d_v}(\text{Compress}_{d_v}(v))$ 
24: cd  $\leftarrow (\mathbf{u}[1][S], \dots, \mathbf{u}[k][S], v[S])$ 
return  $(c = c_1 \| c_2, \text{cd})$ 
```

Verifiable Decapsulation: Confirmation-code Augmented FO

- We formalize confirmation code unpredictability (cUP) for PKE schemes:

limited access to F
intuition: won't accidentally compute

$$\begin{aligned} 03 \quad (c, cd) &\leftarrow \text{Enc}_C^F(\text{pk}, m; r) \\ 04 \quad cd' &\leftarrow \bar{\mathcal{A}}^F(\text{pk}, \text{sk}, c, m, r) \end{aligned}$$

- We introduce a **confirmation-code augmented FO transform** $\text{FOC} = \text{UC} \circ \text{TC}$ [following HHK'17]
 - TC transform: derandomize PKE with confirmation codes
 - UC transform: bind confirmation code into KEM key derivation
- We show: FOC transform of cUP PKE scheme \rightarrow KEM with noticeable incorrectness for faulty implementations

Verifiable Decapsulation without Algorithm Modification

but confirmation codes themselves are

non-intrusive —

$\frac{\text{Decaps}(\text{sk}, c)}{05 \ m' \leftarrow \text{Dec}(\text{sk}, c)}$

06 $(c', \text{cd}') \leftarrow \text{Enc}(\text{pk}, m')$

07 check $c' = c$

08 $K' \leftarrow \text{KDF}(m', \text{pk}, \text{cd}')$

09 **return** K'

intrusive:

requires algorithm modification
= deviation from KEM standards
= problematic for certification



we can post-process confirmation codes via an external wrapper

- treat and formalize **read-only white-box** access to the original KEM implementation
- prove IND-CCA security of original KEM under leakage of confirmation codes
- interesting subtleties: how should an **implicitly rejecting confirmation code** look like?

Summary

(HYBRID) OBFUSCATION

Kemeleon: obfuscate ML-KEM pk/ctxt
– pk even 2.5% smaller



Obfuscated KEM

OEINC: hybrid KEM obfuscation

full versions @ IACR ePrint:

- Kemeleon: ia.cr/2024/1086
<https://datatracker.ietf.org/doc/draft-irtf-cfrg-kemeleon/>
- Hybrid OKEMs: ia.cr/2025/408
- Verifiable Decaps: ia.cr/2025/450

VERIFIABLE DECAPSULATION

functionality conf. code → security

Confirmation-code augmented FO
ML-KEM: 12-20B → detect prob. ~1/3
HQC: 1B → basic tests catch bug
possible w/o algorithm modification

Thank You!
mail@felixguenther.info